

# Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

In [1]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading extenrnal modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

In [3]:

```

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'D:\sufe/人工智能/spring1819_assignment1/assignment1/cs231n/datasets/cifar-10-pyth

    # Cleaning up variables to prevent loading data multiple times (which may cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]
    mask = np.random.choice(num_training, num_dev, replace=False)
    X_dev = X_train[mask]
    y_dev = y_train[mask]

    # Preprocessing: reshape the image data into rows
    X_train = np.reshape(X_train, (X_train.shape[0], -1))
    X_val = np.reshape(X_val, (X_val.shape[0], -1))
    X_test = np.reshape(X_test, (X_test.shape[0], -1))
    X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis = 0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image
    X_dev -= mean_image

    # add bias dimension and transform into columns
    X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
    X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
    X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
    X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

    return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)

```

```
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)
```

```
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

## Softmax Classifier

Your code for this section will all be written inside **cs231n/classifiers/softmax.py**.

In [4]:

```
# First implement the naive softmax loss function with nested loops.
# Open the file cs231n/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs231n.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.330187
sanity check: 2.302585
```

### Inline Question 1

Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

**Your Answer** Since  $W$  is randomly initialized, the score calculated by each class is the same, and the probability after softmax is the same. This is a problem of ten classifications, so the probability of each class is 0.1, and the cross entropy is also  $-\log(0.1)$ .

In [5]:

```

# Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

```

```

numerical: -0.723707 analytic: -0.723707, relative error: 1.932496e-08
numerical: 0.440974 analytic: 0.440974, relative error: 7.816952e-08
numerical: -1.096835 analytic: -1.096835, relative error: 2.675580e-08
numerical: -0.439062 analytic: -0.439062, relative error: 5.926508e-08
numerical: 3.320642 analytic: 3.320642, relative error: 7.963173e-09
numerical: 0.660264 analytic: 0.660264, relative error: 4.134352e-08
numerical: -1.245803 analytic: -1.245803, relative error: 2.178137e-09
numerical: 0.084048 analytic: 0.084048, relative error: 5.414420e-08
numerical: 3.262950 analytic: 3.262950, relative error: 2.610712e-08
numerical: 1.643008 analytic: 1.643008, relative error: 1.507844e-08
numerical: -0.211422 analytic: -0.211422, relative error: 1.257053e-07
numerical: 1.268238 analytic: 1.268239, relative error: 4.137184e-08
numerical: 2.014144 analytic: 2.014143, relative error: 3.760063e-08
numerical: -0.108817 analytic: -0.108817, relative error: 6.907895e-07
numerical: -0.059051 analytic: -0.059051, relative error: 8.750265e-07
numerical: 4.722119 analytic: 4.722119, relative error: 1.830340e-09
numerical: -4.773189 analytic: -4.773189, relative error: 1.008671e-08
numerical: -0.830868 analytic: -0.830868, relative error: 3.034089e-08
numerical: -0.680068 analytic: -0.680068, relative error: 3.127239e-08
numerical: -1.545664 analytic: -1.545664, relative error: 3.764561e-09

```

In [6]:

```
# Now that we have a naive implementation of the softmax loss function and its gradient,  
# implement a vectorized version in softmax_loss_vectorized.  
# The two versions should compute the same results, but the vectorized version should be  
# much faster.  
tic = time.time()  
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)  
toc = time.time()  
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))  
  
from cs231n.classifiers.softmax import softmax_loss_vectorized  
tic = time.time()  
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.000005)  
toc = time.time()  
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))  
  
# As we did for the SVM, we use the Frobenius norm to compare the two versions  
# of the gradient.  
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')  
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))  
print('Gradient difference: %f' % grad_difference)
```

```
naive loss: 2.330187e+00 computed in 0.093727s  
vectorized loss: 2.330187e+00 computed in 0.062493s  
Loss difference: 0.000000  
Gradient difference: 0.000000
```

In [7]:

```

# Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.
from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save   #
# the best trained softmax classifier in best_softmax.                         #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for reg in regularization_strengths:
    for lr in learning_rates:
        svm = Softmax()
        loss_hist = svm.train(X_train, y_train, lr, reg, num_iters=1500)
        y_train_pred = svm.predict(X_train)
        train_accuracy = np.mean(y_train == y_train_pred)
        y_val_pred = svm.predict(X_val)
        val_accuracy = np.mean(y_val == y_val_pred)
        if val_accuracy > best_val:
            best_val = val_accuracy
            best_softmax = svm
        results[(lr, reg)] = train_accuracy, val_accuracy

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)

```

```

lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.352306 val accuracy: 0.359000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.330898 val accuracy: 0.351000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.351878 val accuracy: 0.362000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.329224 val accuracy: 0.339000
best validation accuracy achieved during cross-validation: 0.362000

```

In [8]:

```

# evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

softmax on raw pixels final test set accuracy: 0.366000

```

**Inline Question 2 - True or False**

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

*Your Answer* True

*Your Explanation* Since  $W$  is randomly initialized, the score calculated by each class is the same, and the probability after softmax is the same. This is a very high probability because when the loss function of SVM is calculated, if the newly added test pictures are classified correctly, the loss must be 0. However, for softmax, no matter whether the classification is correct or not, a probability distribution will always be obtained and the cross entropy will be calculated. In other words, softmax's loss will always add a quantity, even a small quantity.

In [9]:

```
# Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```

